

DM207 I/O-Efficient Algorithms and Data Structures

Fall 2015
Project 1

Mikkel Levisen and Jesper Lund

November 24, 2015

Introduction

The goal of this project is to get a feeling of the possible impact of differences in memory access patterns on the running time of programs.

Through this report have the running times been plotted along with an estimate on of cache and ram sizes.

The machines used in this assignment are the ones from the IMADA terminal room and had the following specifications:

Memory type	size	
L1	32	kB
L2	256	kB
L3	4096	kB
RAM	3910084	kB

Since these computers are available to other student at the same time as the tests were running it might have effected the results.

Where appropriate we have added dashed lines indicating the estimated cache and RAM limits. These guides are only estimates and don't consider other running processes.

Task 1




Figure 1: Quicksort vs. Heapsort on arrays of `int`.
(`int`= 4 bytes)

In figure 1 are arrays of sizes between 1000 and 577011370 containing integers sorted with Heapsort and Quicksort. The plot shows the running time of each algorithm relative to the expected running time $O(n \log n)$.

Based on the memory access pattern of the algorithms we are surprised that Heapsort seems to outperform Quicksort while handling arrays which fit in cache. This might be attributed to the lower number of operations which can dominate on instances of small n .

Once the array sizes exceeds the cache limit it clearly show that the spread out memory access of Heapsort is heavily penalized.

Task 2

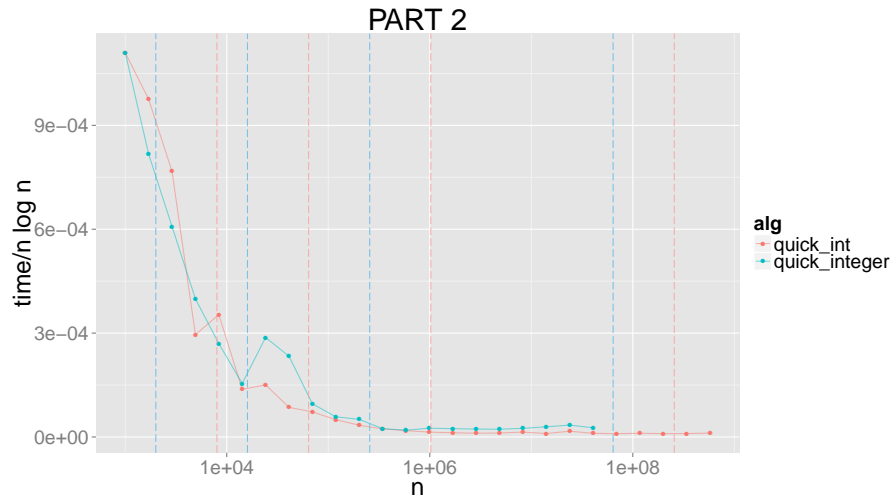


Figure 2: Quicksort `int` vs. Quicksort `Integer`.
(`int`= 4 bytes),`Integer`= 16 bytes)

Because Java uses *pass-by-reference* when dealing with objects with an array of objects simply be pointers to memory locations in the heap which removes Quicksort memory access pattern advantage with sequential data. In figure 2 this is especially visible when the data size reaches L3 cache. In RAM Quicksort still has an advantage because it focuses on distinct partitions of the data rather than the whole set. This leads to a more efficient use of the cache.




Figure 3: Heapsort int vs. Heapsort Integer.
(int= 4 bytes),Integer= 16 bytes)

Heapsort has less of a time penalty since it already has a pseudo random memory access pattern. Therefor it has a similar running time until it reaches the RAM where out algorithm time out.

Task 3




Figure 4: n -cyclic list lookup

Here we clearly see that SATTOLO performs much worse than a standard n -cycle because it has potential for 1 cache miss per memory access. SCAN only has one cache miss per blocksize.

need to load the same blocks in over and over again where SCAN only needs to load a block once.

Task 4





Figure 5: Random array access

As d increases is each block utilized less resulting in an increase in running time. Around $d = 5$ we leave L1 cache, around $d = 13$ we leave L2. A very noticeable increase starting from $d = 17$ indicate that we are querying the RAM regularly. Around $d = 24$ we see that the number of element touched by the algorithm has decreased to a size that fits in the cache and at $d = 27$ we have a running time similar to the initial scan running time.

Task 5

5.1



5.2




Figure 7: Quicksort vs. Heapsort on arrays of `Integer`.
(`Integer`= 16 bytes)

Quicksort still outperforms Heapsort but the effect is less significant because of how Java handles objects. Objects outside of ram can now be stored in different sectors and would dramatically reduce the efficiency of Quicksort.

5.3

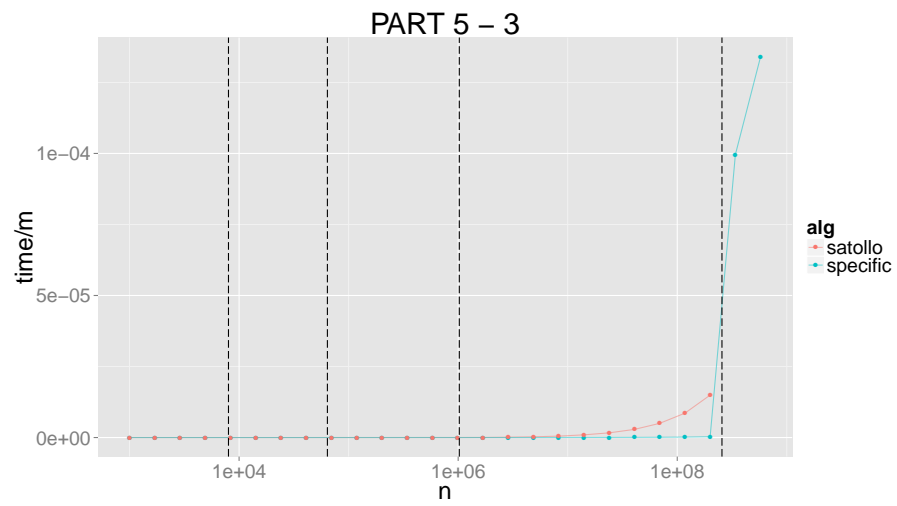


Figure 8: n -cyclic list lookup

Again we see a huge running time penalty from utilizing the disk. Note that SATTOLO timed out when it hit the disk.

5.4




Figure 9: Random array access

Since the array is larger than the RAM + cache is the operation very expensive right from the start, timing out at $i = 3$.